

LEXNOVA AI

Internal Legal AI Platform

Product Vision and Roadmap

From foundation to firm-wide intelligence in eight weeks

Prepared by

Vincent BIROT, Managing Partner

For

Lex Nova Partners — Bangkok, Thailand

Version 1.0 · May 2026

Living document — subject to revision as the build progresses

What is in this document

1.	Executive Summary	p. 3
2.	Vision: Where Lex Nova is going with AI	p. 4
3.	Strategic Value: Build versus Buy	p. 5
4.	Current State: What is operational today	p. 6
5.	Build Roadmap: Eight weeks, five phases	p. 7
6.	Feature Catalogue: Every capability we will ship	p. 8
7.	Investment Summary: Infrastructure, models, build effort	p. 12
8.	Success Metrics: How we will know it is working	p. 13
9.	Risks and Mitigations	p. 14
10.	Next Steps	p. 15

The whole document in 400 words

Lexnova AI is the Lex Nova Partners' internal legal artificial-intelligence platform. It is a private, firm-owned tool, built on the open-source Mike foundation and extended for the firm's specific practice in Bangkok-based corporate, tax, and regulatory law. Lexnova AI replaces ad-hoc lawyer use of Claude.ai and ChatGPT with a unified, secure workspace where documents stay within the firm's perimeter, AI usage is governed and auditable, and the institutional knowledge accumulated across hundreds of past matters becomes searchable, citable, and available to every lawyer at the firm.

As of May 2026, the foundation is operational. Documents upload and are processed. The AI assistant chats with documents and produces high-quality structured output on real Lex Nova matters (smoke testing on a real-estate acquisition matter scored 4 to 5 out of 5 across both summary and precision retrieval tasks). The next eight weeks of focused development, executed autonomously by Claude Code under partner direction, will deliver the operational layer that turns Lexnova AI from a smart document workspace into the firm's daily operating system.

The roadmap covers twelve features in five phases: rebrand and basic UX, operational layer (status, assignee, tasks, time clock, document collaboration), communication integrations (Gmail, Slack), governance and visibility (admin panel, profiles, partner dashboard), and finally the strategic intelligence layer (firm-wide retrieval-augmented generation across closed matters with conflict-wall enforcement).

Total investment is modest: infrastructure runs at USD 65 to 100 per month, model API spend scales with team size and use intensity (USD 200 to 1,500 per month for a 5-lawyer firm), and the build itself consumes roughly USD 1,500 to 3,000 in autonomous Claude Code execution. Compared to Harvey's USD 600 to 1,000 per month for the same team size with a multi-year contract and zero ownership of the codebase, the economics are decisively in favour of building.

The strategic prize is larger than the cost saving. Lex Nova ends this build with a tool that knows the firm's matters, speaks the firm's language, embeds the firm's institutional patterns, and gives every associate from day one the working knowledge of a five-year veteran. No off-the-shelf tool can offer that.

Where Lex Nova is going with AI

Lex Nova Partners aspires to be Thailand's premier corporate, tax, and regulatory practice. In a market where competitors race to AmLaw-100-style scale, Lex Nova has consistently chosen the opposite path: deep, partner-led work, cross-border sophistication, and lasting client relationships.

Lexnova AI is a strategic investment in that vision. Not a productivity tool. Not an experiment. The firm's internal operating layer for legal work over the next decade.

In three years, when a new associate joins Lex Nova, she will have on day one the working knowledge of someone who has been at the firm for five years, because Lexnova AI carries that knowledge for her. When a partner reviews a draft, the AI will already have flagged the same issues a senior associate would have caught. When a client asks 'what is our position on X,' the answer will draw from every matter the firm has handled on X.

Lexnova AI is not a tool the firm uses occasionally for novelty tasks. It is the firm's daily working environment: where matters are managed, where documents are drafted and reviewed, where time is tracked, where research is conducted, where institutional knowledge is preserved and surfaced.

What lawyers do inside Lexnova AI every day, by the end of the build:

- **Open the dashboard** in the morning and see at a glance what is at risk: matters awaiting client more than seven days, overdue tasks, matters approaching their fee cap, the firm's collective capacity for the week.
- **Start a timer** on a matter and capture every billable minute. The timer persists across pages and devices. Manual entries are added retrospectively for time spent off-screen.
- **Chat with the assistant** on any matter, asking real legal questions about real client documents. The AI cites verbatim text, references specific clauses, and where relevant pulls in cross-matter precedent from the firm's archive.
- **Run tabular review** across hundreds of documents in a due-diligence pass. The AI extracts deal terms, governing-law clauses, conditions precedent, and risk flags into a spreadsheet ready for the partner's review.
- **Collaborate on Word drafts** with associates suggesting edits and partners accepting or rejecting them in a single workspace, exporting clean DOCX when ready.
- **See the firm's email and Slack context** on every matter, with the assistant treating linked threads as authoritative facts about what the team and the client have agreed.

This is not an aspirational future. The features that deliver this experience are the twelve briefs in the Lexnova AI fork's *docs/design-briefs/* folder, each ready for autonomous implementation by Claude Code. The build sequence below moves the firm from the current foundation to this vision in eight weeks of calendar time.

Build versus buy, and why this is the right move

The off-the-shelf alternatives

The legal AI market in 2026 is dominated by Harvey (raised at a USD 5B+ valuation), Legora, and a thickening second tier. They are well-engineered, but they are built for AmLaw-100 firms practicing US and UK law on US and UK documents. Their pricing model assumes 50-plus lawyers. Their training data assumes English. Their matter taxonomies assume the practice patterns of large American firms.

Lex Nova is none of those things. The firm practices Thai law, in Bangkok, on bilingual documents, with cross-border European and ASEAN clients, in a boutique structure where the partner reviews every matter. An off-the-shelf legal AI tool will under-serve the firm in exactly the places where it most needs help.

Why building on Mike is the right path

Mike (mikeoss.com) is a working open-source legal-AI platform released in April 2026. It provides the foundation that would have taken Lex Nova nine months to build internally: AI chat with documents, tabular review for batch extraction, document version control with track changes, a mature React + Next.js frontend, an Express + Supabase backend, AGPL-3.0 licensed.

Building on Mike gives Lex Nova:

- **A nine-month head start.** The hard parts of legal AI (document chunking, citation grounding, accept-reject editing, multi-model routing) are already solved.
- **Full ownership of the codebase** after fork. No vendor lock-in. Lex Nova's modifications belong to Lex Nova.
- **Data residency under firm control.** Every document stays in the firm's own Supabase project and Cloudflare R2 bucket. Privilege protection by infrastructure, not by trust.
- **Customisation in directions Harvey will never go.** Thai bilingual handling. Thai-specific OCR. BOI workflow templates. Lex Nova's matter taxonomy. The firm's hourly rates.
- **Cross-matter institutional memory.** Every past matter becomes searchable context for the AI. New associates inherit the firm's accumulated wisdom on day one.
- **Cost economics that scale with the firm, not against it.** Marginal cost per lawyer added is the API spend, not a subscription seat.

The cost comparison

Configuration	Harvey or Legora	Lexnova AI (this build)
5 lawyers, monthly	USD 600 - 1,000 + multi-year contract	USD 250 - 1,500 full ownership
15 lawyers, monthly	USD 1,800 - 3,000 + multi-year contract	USD 700 - 3,500 full ownership
Year-1 build cost	0 (subscription only)	USD 1,500 - 3,000 (autonomous Claude Code)
Codebase ownership	None	Full (private fork)
Thai-specific capability	Limited / generic	Built for Thai practice
Cross-matter institutional memory	Limited; firm-wide RAG is enterprise-tier only	Built into the foundation (Phase 5)

What is operational at Lex Nova as of May 2026

Lexnova AI is already running. Not as a future project; as a working system on Lex Nova's infrastructure. The foundational layer is complete, validated, and ready to be extended.

Infrastructure operational

- Self-hosted on Lex Nova-controlled infrastructure: Supabase project for the database and authentication, Cloudflare R2 bucket for document storage, Anthropic API for the chat assistant, all keys held by the firm.
- Frontend (Next.js) and backend (Express) running locally on the partner's Mac, ready to be deployed to a Lex Nova subdomain (*ai.lexnovapartners.com*) when the team rollout phase begins.
- Private fork of Mike at github.com/vincentbirot/lexnova-mike, with the upstream tracked for future security patches.
- Claude Code CLI installed and validated. The autonomous build loop has been demonstrated end-to-end with the first feature (rename of Project to Matter throughout the user interface) shipped as pull request #1.

Capabilities working today

- **Matter creation and document upload:** PDF and DOCX support. LibreOffice handles the DOCX-to-PDF conversion.
- **AI chat with documents:** lawyers ask questions, the assistant cites verbatim text from documents with page references. Smoke tested on a real Lex Nova matter (real-estate acquisition for Mr. Mathias Barbé) scoring 4 to 4.5 out of 5 on structured summary and 4.5 to 5 out of 5 on precision retrieval of conditions precedent.
- **Tabular review:** extract structured fields across multiple documents in batch.
- **Document version control:** track changes at the cell level, with accept and reject pending in the schema (UI surface comes in Brief 05).
- **Multi-model support:** Claude (default for chat) and Gemini (default for tabular review) with per-user override.

What is not yet there (and is what this roadmap delivers)

- Branding (still 'Mike' in places, will become 'Lexnova AI' as the first feature shipped).
- Operational layer: matter status, assignees, tasks, time tracking.
- Communication: Gmail and Slack integration.
- Governance: admin panel, audit log, usage dashboards.

- Strategic intelligence: cross-matter retrieval (the firm-wide RAG memory).

Eight weeks, five phases, twelve features

The build is sequenced to deliver value at every step. Phase 1 ships in days, gives the team a branded, queue-enabled tool. Phase 2 ships in two weeks, gives the team operational management capability. The remaining phases compound on top.

Phase	Theme	Sprint	Effort	Outcome
PHASE 0	Foundation	Complete	5 days (done)	Stack stood up, documents working, AI chat functional, fork created, smoke testing passed.
PHASE 1	Branding and basic UX	Sprint 0-1	1 - 2 days	Brief 00 (rebrand to Lexnova AI), Brief 01 (concurrent request queue, Claude.ai-style). Outcome: firm-branded tool with daily-use UX parity to Claude.ai.
PHASE 2	Operational layer	Sprint 1-2	1 - 2 weeks	Brief 02 (matter status and assignee), Brief 03 (tasks per matter), Brief 04 (time clock and timesheet), Brief 05 (DOCX collaboration light). Outcome: lawyers manage matters, tasks, and time inside the platform.
PHASE 3	Communication integrations	Sprint 3-4	2 - 3 weeks	Brief 06 (Gmail integration, with Google OAuth verification calendar time of 5 to 7 business days), Brief 07 (Slack integration). Outcome: email and Slack feed matter context; the assistant has the full picture.
PHASE 4	Governance and visibility	Sprint 5-6	1 - 2 weeks	Brief 08 (admin panel), Brief 09 (user profile and templates), Brief 10 (partner dashboard). Outcome: firm leadership has end-to-end oversight; users have a personalised experience.
PHASE 5	Strategic intelligence	Sprint 7+	1 - 2 weeks	Brief 11 (firm-wide RAG memory layer with conflict walls, pgvector embeddings, anonymised cross-matter retrieval). Outcome: every past Lex Nova matter is searchable context for the AI; new associates get day-one institutional knowledge.

Total calendar time: roughly eight weeks, of which four to five weeks are actual focused build effort. The remaining time is third-party verification waits (Google OAuth) and the partner's review cycles between feature ships.

Every capability we will ship

Each feature below is documented in full detail in the corresponding design brief at *docs/design-briefs/* in the Lex Nova fork. The summaries here are the partner-level pitch.

BRIEF 00 Audience All users Effort 4 - 6 hours	Platform Rebrand Mike → Lexnova AI throughout the user interface, browser tabs, login screens, error pages, and source code. Every visible mention of 'Mike' becomes 'Lexnova AI'. Lex Nova's logo replaces Mike's chrysanthemum mark. Source-code attribution to upstream Mike is preserved in a clearly visible footer (an AGPL licence obligation, not just a courtesy).
BRIEF 01 Audience All users Effort 1 - 2 days	Concurrent Request Queue Lawyers can fire multiple AI questions in quick succession without the input blocking. Responses stream concurrently. A queue indicator shows in-flight requests. Multi-tab and multi-matter parallelism work cleanly. This brings the user experience to parity with Claude.ai and removes the daily-use blocker that frustrates the team today.
BRIEF 02 Audience All users Effort 1 day	Matter Status and Assignee Every matter gets a lifecycle status (active, awaiting client, in review, billed, closed) and a single primary assignee. Both visible everywhere matters appear. Both filterable. The default matter list shows only active matters; closed and billed are hidden until explicitly searched. Foundation for everything operational that follows.
BRIEF 03 Audience All users Effort 1 - 2 days	Tasks per Matter Task list on every matter, with due dates and assignees. A 'My Tasks' sidebar surfaces what each lawyer owes across all their matters, grouped by overdue, today, this week, later. A full-page /tasks dashboard for cross-matter task management. Tasks are how a partner sees who is doing what.
BRIEF 04 Audience All users Effort 2 days	Time Clock and Timesheet Per-matter timer with start, stop, and manual retrospective entries. Personal weekly timesheet visible at /time. Per-matter monthly billable rollup ready for invoicing. The single most valuable feature for revenue capture: industry data shows lawyers lose 15 to 25 percent of billable time without integrated tracking.

BRIEF 05 Audience Drafters and reviewers Effort 2 days	DOCX Collaboration (Light) Multi-user accept and reject of edits on Word documents. Associates propose edits in a Lexnova AI viewer, partners review and accept or reject each one, the result exports as clean DOCX. No real-time co-editing in this version (separate workstream). Replaces the email-and-attachment loop the firm uses today.
BRIEF 06 Audience All users Effort 3 - 4 days build + 5 to 7 days Google verification wait	Gmail Integration Per-user Gmail connection via OAuth. Threads link to matters automatically by client domain or manually. The matter detail page surfaces an Email tab. Sent replies use the lawyer's Gmail account. The /inbox view shows matter-relevant threads grouped by needs-reply and awaiting-client. The assistant uses linked threads as context.
BRIEF 07 Audience All users Effort 2 - 3 days	Slack Integration Workspace-level OAuth for the Lex Nova Slack workspace. Each matter optionally maps to a Slack channel (auto-created at matter creation). The matter Slack tab shows recent channel messages. Notifications post to channels on matter and task events. The assistant uses linked Slack messages as context.
BRIEF 08 Audience Admin / managing partners Effort 2 - 3 days	Admin Panel User management (invite, role change, disable), integration management (Anthropic, Gemini, Gmail, Slack, R2), firm-wide template library, audit log, billing and usage dashboards, firm-settings. Reserved for admin role. The governance layer.
BRIEF 09 Audience All users Effort 1 - 2 days	User Profile and Templates Per-lawyer profile with model preferences, billing rate, email signature, personal task templates, personal workflow templates, notification settings, connected integrations. Onboarding wizard for new associates joining the firm. Each lawyer makes Lexnova AI their own.

BRIEF 10 Audience Partners and admin Effort 2 days	Partner Dashboard Personalised /dashboard showing matters at risk, today's tasks, recent activity, team capacity (admin view), awaiting-client matters, and AI usage. The morning landing page for partners. Answers the daily question 'where is everything and what needs my attention'.
BRIEF 11 Audience All users Effort 1 - 2 weeks	Firm-Wide RAG Memory Pgvector embeddings of every document across every matter. The assistant retrieves cross-matter precedent on every chat turn (subject to access rules and conflict walls). A /search route for explicit firm-wide search. Conflict-wall enforcement protects privilege between matters that conflict. The single biggest strategic feature in the build.

What it costs to build and run Lexnova AI

Infrastructure (recurring monthly, production)

Item	Description	Monthly cost
Supabase Pro	Database, authentication, storage tier with daily backups and point-in-time recovery	USD 25
Cloudflare R2	Object storage for documents, zero egress fees, S3-compatible API	USD 3 - 10
Vercel Pro	Frontend hosting, commercial-licence tier, password protection, custom domain	USD 20
Backend hosting	Railway or Fly.io for Express + LibreOffice container	USD 15 - 30
Resend (email)	Transactional email for notifications and invitations	USD 0 - 15
Subtotal infrastructure		USD 65 - 100

Model API spend (variable, scales with team size and use)

Configuration	Anthropic	Gemini	Per month total
5 lawyers, light use	USD 150 - 500	USD 25 - 100	USD 200 - 600
5 lawyers, heavy DD use	USD 500 - 1,250	USD 100 - 250	USD 600 - 1,500
15 lawyers, mixed use	USD 600 - 2,500	USD 100 - 500	USD 700 - 3,000

One-time build cost

The features are implemented by Claude Code running autonomously on the partner's Mac (or in a cloud runner) from the briefs in *docs/design-briefs/*. Total Anthropic API consumption for the build, across all twelve features and all review cycles, is estimated at **USD 1,500 to 3,000**. Partner review and direction time across the eight-week build is roughly **30 to 50 hours**, distributed across the period.

How we will know it is working

Lexnova AI is a strategic investment, not a productivity experiment. The metrics that matter are the ones that show whether the firm is using it daily, whether it is making the firm faster and more accurate, and whether it is becoming the institutional memory layer the vision requires.

Adoption (Phase 1 - 2)

- Daily active users: target 80 percent of the firm logged in on a given working day by end of Sprint 2.
- Matters with at least one document uploaded: target 90 percent of new matters by end of Sprint 2.
- Time entries created: target average 5 hours of tracked time per lawyer per working day by end of Sprint 2 (versus the firm's current rough estimate of 3 to 4 hours captured manually).

Quality (continuous)

- Citation accuracy on legal documents: target 90 percent verbatim accuracy and 95 percent page-reference accuracy on a representative sample of Lex Nova matter documents (audited monthly).
- Hallucination rate: target less than 2 percent on factual claims sampled monthly. Any matter where the assistant produces a hallucinated Thai legal citation triggers a remediation review.
- Thai-language handling: target equal-quality output on Thai versus English versions of bilingual contracts (formally evaluated quarterly).

Operational impact (Phase 2 - 4)

- Matter throughput: target 10 to 15 percent more matters handled per partner per quarter by end of Phase 4.
- Awaiting-client follow-up time: target 50 percent reduction in matters awaiting client response more than 7 days, by end of Phase 4 (driven by the dashboard surfacing the gap).
- Billable-time capture rate: target 95 percent of working time captured as time entries (industry benchmark for top firms is 92 - 96 percent).

Strategic impact (Phase 5 onward)

- New-associate productivity: target a new associate billing 60 percent of partner-equivalent productivity within 60 days of joining (versus the industry average of 6 to 12 months for partner-equivalent productivity).
- Cross-matter precedent surfacing: at least 40 percent of legal-research questions surface relevant cross-matter precedent automatically by end of Phase 5.

- Client retention: track whether AI-augmented service quality correlates with the firm's repeat-engagement rate over the year following Phase 5 launch.

What could go wrong and how we are protecting against it

Privilege breach	A document or chat from one matter leaks to a user who should not see it (e.g., a conflict-walled matter).	<i>Conflict-wall mechanism in Phase 5 enforces matter-level access. All documents stored in firm-controlled Supabase + R2. The assistant's retrieval is filtered through the access rules on every query. Quarterly access audit against the audit log.</i>
Hallucinated Thai legal citation	The assistant cites a non-existent Thai statute or misstates a real one.	<i>Citation grounding in source documents (the assistant cites verbatim text, not its training data). Monthly hallucination audit on a sample of outputs. Lawyer-verification expectation set culturally: 'AI-assisted, lawyer-verified' is the standard for any client-facing output.</i>
AI provider outage	Anthropic or Gemini goes down for a sustained period.	<i>Multi-provider support already built into Mike (Claude, Gemini, OpenRouter). Falls back to alternate provider on outage. Critical matters can be handled in pre-AI workflow during downtime.</i>
Cost runaway	A lawyer or batch operation consumes far more API spend than expected.	<i>Per-key spend caps on Anthropic and Gemini. Budget alerts at 50, 80, 100 percent thresholds. Per-user usage visible in admin panel (Brief 08). Tabular review explicit batch limits to prevent accidental large runs.</i>
Google OAuth verification delay	Gmail integration blocked while waiting for Google's verification of sensitive scopes.	<i>Submit verification on Day 1 of Sprint 3. Build the integration UI in parallel during the wait. Integration goes live to all users when verification completes. Until then, only the submitting partner can use it.</i>
Mike upstream divergence	The open-source Mike project changes in ways that conflict with our fork.	<i>Maintain the upstream remote and pull periodically. Any conflict is a manageable code merge handled by Claude Code. Lex Nova's fork is the source of truth for the firm; upstream changes are evaluated for inclusion case by case.</i>

Adoption stall	Lawyers do not change their daily habits to use Lexnova AI.	<i>Sprint 1 (queue + branding) addresses the most-cited frustration. Sprint 2 (operational layer) makes Lexnova AI the only place to log time and tasks. Sprint 4 (Slack integration) embeds Lexnova AI in the team's existing communication. Adoption is engineered, not hoped for.</i>
PDPA / data-privacy issue	Client data crosses borders or is processed in ways that violate Thai data-protection rules.	<i>All documents stored in firm-controlled Supabase + R2 with region selection. AI providers (Anthropic, Gemini) accessed under their data-processing agreements with no-training-on-data terms. Per-matter 'exclude from RAG' setting in Phase 5 for clients who require absolute confidentiality. PDPA assessment formally reviewed before Phase 5 launch.</i>

What happens after this document is shared

Immediate (this week)

- Partners read this document and the design briefs in *docs/design-briefs/*. Push back on anything that does not match Lex Nova's intent.
- Provide brand assets for the rebrand (Brief 00): logo SVG, mark for sidebar collapsed state, brand color hex values, optional tagline.
- Decide on cloud-deployment timing: Phase 1 ships to localhost first; the team rollout phase begins when Phase 2 is in place and the Lex Nova subdomain is wired up.

Sprint 0 - 1 (next 1 - 2 working days)

- Run Brief 00 (rebrand) through Claude Code Design (or v0 / Lovable) to produce visual assets and copy strings, then through Claude Code to implement.
- Run Brief 01 (queue) through the same loop. Review and merge.
- Outcome: branded, queue-enabled Lexnova AI ready for daily use by the partner.

Sprint 2 (week 2)

- Operational layer (Briefs 02 - 05). Three Claude Code sessions in parallel produce three pull requests.
- Outcome: status, assignee, tasks, time clock, and DOCX collaboration shipped.
- Cloud deployment to *ai.lexnovapartners.com* initiated (Vercel for frontend, Railway or Fly for backend, DNS via Cloudflare).

Sprints 3 - 7 (weeks 3 - 8)

- Communication integrations (Brief 06 Gmail submitted to Google for OAuth verification on Day 1, Brief 07 Slack built in parallel).
- Governance layer (Briefs 08 - 10).
- Strategic intelligence layer (Brief 11 RAG memory).
- Outcome by end of Sprint 7: full Lexnova AI vision operational across the firm.

After Phase 5

- Continuous improvement based on real lawyer feedback.
- Optional submission of generally-useful fixes upstream to the Mike project (the AGPL-3.0 licence does not require this, but it is good citizenship and benefits the broader legal-AI

ecosystem).

- Quarterly review of metrics in Section 8. Adjust roadmap based on what the firm actually values most after using Lexnova AI for 90 days.

This is a living document. Versions and revisions are tracked in the Lex Nova fork at github.com/vincentbirot/lexnova-mike/blob/main/docs/Lexnova-AI-Product-Vision.pdf. The technical design briefs sit alongside it at [docs/design-briefs/](#).